

# Quantium Virtual Internship - Retail Strategy and Analytics - Task 1

Sarita

2024-05-13

## Load required libraries

```
library(data.table)
library(ggplot2)
library(ggmosaic)
library(readr)
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.4      v stringr  1.5.1
## v forcats    1.0.0      v tibble   3.2.1
## v lubridate  1.9.3      v tidyr    1.3.1
## v purrr      1.0.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::between()      masks data.table::between()
## x dplyr::filter()       masks stats::filter()
## x dplyr::first()        masks data.table::first()
## x lubridate::hour()     masks data.table::hour()
## x lubridate::isoweek()  masks data.table::isoweek()
## x dplyr::lag()          masks stats::lag()
## x dplyr::last()         masks data.table::last()
## x lubridate::mday()     masks data.table::mday()
## x lubridate::minute()   masks data.table::minute()
## x lubridate::month()    masks data.table::month()
## x lubridate::quarter()  masks data.table::quarter()
## x lubridate::second()   masks data.table::second()
## x purrr::transpose()    masks data.table::transpose()
## x lubridate::wday()     masks data.table::wday()
## x lubridate::week()     masks data.table::week()
## x lubridate::yday()     masks data.table::yday()
## x lubridate::year()     masks data.table::year()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

## import csv

```
filePath <- ""
transactionData <- fread(paste0(filePath,"QVI_transaction_data.csv"))
customerData <- fread(paste0(filePath,"QVI_purchase_behaviour.csv"))
```

```
head(transactionData)
```

```
##      DATE STORE_NBR LYLTY_CARD_NBR TXN_ID PROD_NBR
##      <int>      <int>          <int> <int>   <int>
## 1: 43390         1           1000     1      5
## 2: 43599         1           1307    348     66
## 3: 43605         1           1343    383     61
## 4: 43329         2           2373    974     69
## 5: 43330         2           2426   1038    108
## 6: 43604         4           4074   2982     57
##                                     PROD_NAME PROD_QTY TOT_SALES
##                                     <char>    <int>    <num>
## 1:   Natural Chip          Compny SeaSalt175g         2        6.0
## 2:                CCs Nacho Cheese    175g         3        6.3
## 3:   Smiths Crinkle Cut  Chips Chicken 170g         2        2.9
## 4:   Smiths Chip Thinly  S/Cream&Onion 175g         5       15.0
## 5: Kettle Tortilla ChpsHny&Jlpno Chili 150g         3       13.8
## 6: Old El Paso Salsa    Dip Tomato Mild 300g         1        5.1
```

```
head(customerData)
```

```
##      LYLTY_CARD_NBR          LIFESTAGE PREMIUM_CUSTOMER
##      <int>          <char>          <char>
## 1:         1000  YOUNG SINGLES/COUPLES      Premium
## 2:         1002  YOUNG SINGLES/COUPLES      Mainstream
## 3:         1003          YOUNG FAMILIES      Budget
## 4:         1004  OLDER SINGLES/COUPLES      Mainstream
## 5:         1005  MIDAGE SINGLES/COUPLES      Mainstream
## 6:         1007  YOUNG SINGLES/COUPLES      Budget
```

## Exploratory data analysis

### TransactionData

```
str(transactionData)
```

```
## Classes 'data.table' and 'data.frame': 264836 obs. of 8 variables:
## $ DATE : int 43390 43599 43605 43329 43330 43604 43601 43601 43332 43330 ...
## $ STORE_NBR : int 1 1 1 2 2 4 4 4 5 7 ...
## $ LYLTY_CARD_NBR: int 1000 1307 1343 2373 2426 4074 4149 4196 5026 7150 ...
## $ TXN_ID : int 1 348 383 974 1038 2982 3333 3539 4525 6900 ...
## $ PROD_NBR : int 5 66 61 69 108 57 16 24 42 52 ...
## $ PROD_NAME : chr "Natural Chip Compny SeaSalt175g" "CCs Nacho Cheese 175g"
## "Smiths Crinkle Cut Chips Chicken 170g" "Smiths Chip Thinly S/Cream&Onion 175g"
## ...
## $ PROD_QTY : int 2 3 2 5 3 1 1 1 1 2 ...
## $ TOT_SALES : num 6 6.3 2.9 15 13.8 5.1 5.7 3.6 3.9 7.2 ...
## - attr(*, ".internal.selfref")=<externalptr>
```

```
transactionData %>%
  summarise_all(class) %>%
  gather(variale, class)
```

```
##          variable      class
## 1          DATE      integer
## 2      STORE_NBR      integer
## 3 LYLTY_CARD_NBR      integer
## 4          TXN_ID      integer
## 5      PROD_NBR      integer
## 6      PROD_NAME character
## 7      PROD_QTY      integer
## 8      TOT_SALES      numeric
```

## CustomerData

```
str(customerData)
```

```
## Classes 'data.table' and 'data.frame': 72637 obs. of 3 variables:
## $ LYLTY_CARD_NBR : int 1000 1002 1003 1004 1005 1007 1009 1010 1011 1012 ...
## $ LIFESTAGE : chr "YOUNG SINGLES/COUPLES" "YOUNG SINGLES/COUPLES" "YOUNG
FAMILIES" "OLDER SINGLES/COUPLES" ...
## $ PREMIUM_CUSTOMER: chr "Premium" "Mainstream" "Budget" "Mainstream" ...
## - attr(*, ".internal.selfref")=<externalptr>
```

```
customerData %>%
  summarise_all(class) %>%
  gather(variable, class)
```

```
##          variable      class
## 1  LYLTY_CARD_NBR      integer
## 2      LIFESTAGE character
## 3 PREMIUM_CUSTOMER character
```

## Examining TransactionData data

From exploring the dataset I found that the date column is in an integer format, so the first step is changing the date format.

### Converting Date Format

```
transactionData$DATE <- as.Date(transactionData$DATE, origin = "1899-12-30")
head(transactionData)
```

```
##          DATE STORE_NBR LYLTY_CARD_NBR TXN_ID PROD_NBR
##          <Date>      <int>          <int> <int>      <int>
## 1: 2018-10-17          1          1000      1          5
## 2: 2019-05-14          1          1307     348         66
## 3: 2019-05-20          1          1343     383         61
## 4: 2018-08-17          2          2373     974         69
## 5: 2018-08-18          2          2426    1038        108
## 6: 2019-05-19          4          4074     2982         57
##                                     PROD_NAME PROD_QTY TOT_SALES
##                                     <char>      <int>      <num>
## 1: Natural Chip          Compny SeaSalt175g          2          6.0
## 2:                CCs Nacho Cheese    175g          3          6.3
## 3: Smiths Crinkle Cut Chips Chicken 170g          2          2.9
```

```
## 4:   Smiths Chip Thinly  S/Cream&Onion 175g      5      15.0
## 5: Kettle Tortilla ChpsHny&Jlpno Chili 150g      3      13.8
## 6: Old El Paso Salsa   Dip Tomato Mild 300g      1       5.1
```

## Examining PROD\_NAME

```
transactionData %>% count(PROD_NAME)
```

```
##              PROD_NAME      n
##              <char> <int>
##    1:      Burger Rings 220g 1564
##    2:      CCs Nacho Cheese 175g 1498
##    3:      CCs Original 175g 1514
##    4:      CCs Tasty Cheese 175g 1539
##    5:      Cheetos Chs & Bacon Balls 190g 1479
## ---
## 110: WW Sour Cream &OnionStacked Chips 160g 1483
## 111: WW Supreme Cheese   Corn Chips 200g 1509
## 112:      Woolworths Cheese   Rings 190g 1516
## 113:      Woolworths Medium   Salsa 300g 1430
## 114:      Woolworths Mild     Salsa 300g 1491
```

There are 114 types of product but we are only interested in the potato chips, so we would like to keep only the data of potato ships and discard other by summarising the individual words in the product name.

```
productWords <- data.table(unlist(strsplit(unique(transactionData$PROD_NAME), " ")))
setnames(productWords, "words")
```

## Removing digits

```
productWords <- productWords[grepl("[[:digit:]]", words) == FALSE, ]
```

## Removing special characters

```
productWords <- productWords[grepl("[[:punct:]]", words) == FALSE, ]
```

## Sorting by frequency

```
productWords[, .N, words][order(N, decreasing = TRUE)]
```

```
##      words      N
##      <char> <int>
##    1:      234
##    2:    Chips   21
##    3:   Smiths   16
##    4:  Crinkle   14
##    5:     Cut    14
## ---
## 165:     Rst     1
## 166:    Pork     1
## 167:   Belly     1
## 168:     Pc      1
```

## 169: Bolognese 1

## Remove the salsa product

```
transactionData[, SALSA := grepl("salsa", tolower(PROD_NAME))]  
transactionData <- transactionData[SALSA == FALSE, ][, SALSA := NULL]
```

## Summarizing the data

```
summary(transactionData)
```

```
##      DATE      STORE_NBR  LYLTY_CARD_NBR  TXN_ID  
## Min.   :2018-07-01  Min.   : 1.0  Min.   : 1000  Min.   : 1  
## 1st Qu.:2018-09-30  1st Qu.: 70.0  1st Qu.: 70015  1st Qu.: 67569  
## Median :2018-12-30  Median :130.0  Median : 130367  Median : 135183  
## Mean   :2018-12-30  Mean   :135.1  Mean   : 135531  Mean   : 135131  
## 3rd Qu.:2019-03-31  3rd Qu.:203.0  3rd Qu.: 203084  3rd Qu.: 202654  
## Max.   :2019-06-30  Max.   :272.0  Max.   :2373711  Max.   :2415841  
##      PROD_NBR      PROD_NAME      PROD_QTY      TOT_SALES  
## Min.   : 1.00  Length:246742  Min.   : 1.000  Min.   : 1.700  
## 1st Qu.: 26.00  Class :character  1st Qu.: 2.000  1st Qu.: 5.800  
## Median : 53.00  Mode  :character  Median : 2.000  Median : 7.400  
## Mean   : 56.35                      Mean   : 1.908  Mean   : 7.321  
## 3rd Qu.: 87.00                      3rd Qu.: 2.000  3rd Qu.: 8.800  
## Max.   :114.00                      Max.   :200.000  Max.   :650.000
```

There are no nulls in the columns but product quantity appears to have an outlier which case where 200 packets of chips are bought in one transaction.

## Filter the dataset to find the outlier

```
transactionData[PROD_QTY == 200, ]
```

```
##      DATE STORE_NBR LYLTY_CARD_NBR TXN_ID PROD_NBR  
##      <Date>      <int>          <int> <int>      <int>  
## 1: 2018-08-19      226          226000 226201        4  
## 2: 2019-05-20      226          226000 226210        4  
##      PROD_NAME PROD_QTY TOT_SALES  
##      <char>      <int>      <num>  
## 1: Dorito Corn Chp Supreme 380g      200      650  
## 2: Dorito Corn Chp Supreme 380g      200      650
```

There are two transactions where 200 packets of chips are bought in one transaction and both of these transactions were by the same customer.

## Examining other transactions that customer made

```
transactionData[LYLTY_CARD_NBR == 226000, ]
```

```
##      DATE STORE_NBR LYLTY_CARD_NBR TXN_ID PROD_NBR  
##      <Date>      <int>          <int> <int>      <int>  
## 1: 2018-08-19      226          226000 226201        4  
## 2: 2019-05-20      226          226000 226210        4  
##      PROD_NAME PROD_QTY TOT_SALES  
##      <char>      <int>      <num>  
## 1: Dorito Corn Chp Supreme 380g      200      650
```

```
## 2: Dorito Corn Chp      Supreme 380g      200      650
```

This customer has only had the two transactions over the year and is not an ordinary retail customer. The customer might be buying chips for commercial purposes instead. We'll remove this loyalty card number from further analysis.

Filter out the outlier

```
transactionData <- transactionData[LYLTY_CARD_NBR != 226000, ]
```

Re-examine transaction data

```
summary(transactionData)
```

```
##      DATE      STORE_NBR  LYLTY_CARD_NBR      TXN_ID
## Min.   :2018-07-01  Min.   :  1.0  Min.   :  1000  Min.   :    1
## 1st Qu.:2018-09-30  1st Qu.: 70.0  1st Qu.:  70015  1st Qu.: 67569
## Median :2018-12-30  Median :130.0  Median : 130367  Median : 135182
## Mean   :2018-12-30  Mean   :135.1  Mean   : 135530  Mean   : 135130
## 3rd Qu.:2019-03-31  3rd Qu.:203.0  3rd Qu.: 203083  3rd Qu.: 202652
## Max.   :2019-06-30  Max.   :272.0  Max.   :2373711  Max.   :2415841
##  PROD_NBR  PROD_NAME  PROD_QTY  TOT_SALES
## Min.   :  1.00  Length:246740  Min.   :1.000  Min.   :  1.700
## 1st Qu.: 26.00  Class :character  1st Qu.:2.000  1st Qu.:  5.800
## Median : 53.00  Mode  :character  Median :2.000  Median :  7.400
## Mean   : 56.35  Mean   :1.906  Mean   :  7.316
## 3rd Qu.: 87.00  3rd Qu.:2.000  3rd Qu.:  8.800
## Max.   :114.00  Max.   :5.000  Max.   :29.500
```

Now, let's examine the number of transaction lines over time to find obvious data issues such as missing data.

## Finding obvious data issue

Counting the number of transactions by date

```
transactionData[, .N, by = DATE]
```

```
##      DATE      N
##      <Date> <int>
## 1: 2018-10-17  682
## 2: 2019-05-14  705
## 3: 2019-05-20  707
## 4: 2018-08-17  663
## 5: 2018-08-18  683
## ---
## 360: 2018-12-08  622
## 361: 2019-01-30  689
## 362: 2019-02-09  671
## 363: 2018-08-31  658
## 364: 2019-02-12  684
```

There's only 364 rows, meaning only 364 dates which indicates a missing date. Next, I will create a sequence of dates from 1 Jul 2018 to 30 Jun 2019 and use this to create a chart of number of transactions over time to find the missing date.



### Filling the missing day

Creating a sequence of dates and join this the count of transactions by date to fill in the missing day


```
allDates <- data.table(seq(as.Date("2018/07/01"), as.Date("2019/06/30"), by = "day"))
setnames(allDates, "DATE")
transactions_by_day <- merge(allDates, transactionData[, .N, by = DATE], all.x = TRUE)
```

### Setting plot themes to format graphs

```
theme_set(theme_bw())
theme_update(plot.title = element_text(hjust = 0.5))
```

### Plotting transactions over time

```
ggplot(transactions_by_day, aes(x = DATE, y = N)) +
  geom_line(col = "red") +
  labs(x = "Day", y = "Number of transactions", title = "Transactions over time") +
  scale_x_date(breaks = "1 month") +
  theme(axis.text.x = element_text(angle = 90, vjust = 0.5))
```




We can see that there is an increase in purchases in December and a break in late December. Let's zoom in on this.



Filtering to December and look at individual days

```
ggplot(transactions_by_day[month(DATE)==12, ], aes(x = DATE, y = N)) +  
  geom_line(col = "blue") +  
  labs(x = "Day", y = "Number of transactions", title = "Transactions over time") +  
  scale_x_date(breaks = "1 day") +  
  theme(axis.text.x = element_text(angle = 90, vjust = 0.5))
```



We can see that the increase in sales occurs in the lead-up to Christmas and that there are zero sales on Christmas day itself. This is due to shops being closed on Christmas day.

Now, the data no longer has outliers and any other missing value. So we can move on to creating other features. I will start with pack size from PROD\_NAME.

## Examining pack size

```
transactionData[, PACK_SIZE := parse_number(PROD_NAME)]  
transactionData[, .N, PACK_SIZE][order(PACK_SIZE)]
```

```
##   PACK_SIZE    N  
##   <num> <int>  
## 1:      70 1507  
## 2:      90 3008  
## 3:     110 22387  
## 4:     125 1454  
## 5:     134 25102  
## 6:     135 3257
```

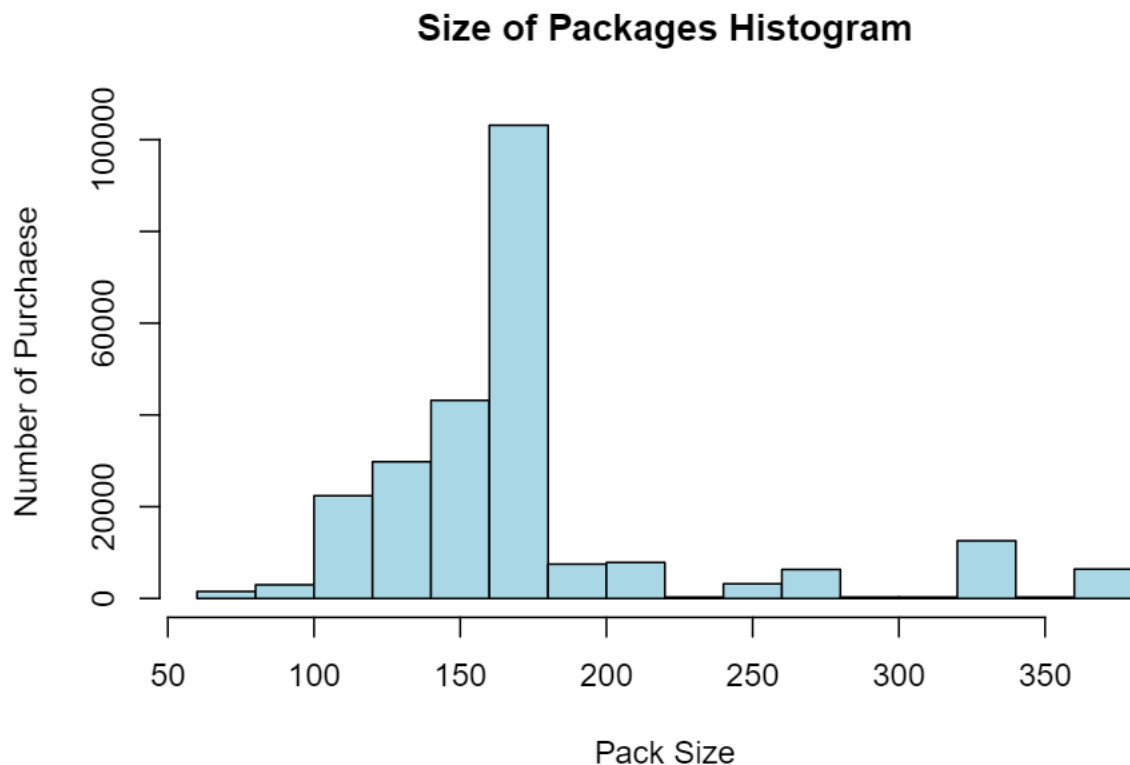
```
## 7:      150 40203
## 8:      160  2970
## 9:      165 15297
## 10:     170 19983
## 11:     175 66390
## 12:     180  1468
## 13:     190  2995
## 14:     200  4473
## 15:     210  6272
## 16:     220  1564
## 17:     250  3169
## 18:     270  6285
## 19:     330 12540
## 20:     380  6416
##      PACK_SIZE      N
```

The largest size is 380g and the smallest size is 70g - seems sensible.

Next, we will plot a histogram of PACK\_SIZE since we know that it's a categorical variable and not a continuous variable even though it's numeric.

#### Plotting Histogram of PACK\_SIZE

```
options(scipen = 999)
hist(transactionData$PACK_SIZE,
      col = "lightblue",
      main = "Size of Packages Histogram",
      xlab = "Pack Size",
      ylab = "Number of Purchase")
```



The histogram looks reasonable. We found that the packs of size 150-200 was purchased the most.

## Brand

To creating BRAND, we will use the first word in PROD\_NAME to work out the brand name.

```
transactionData[, BRAND := word(transactionData[, PROD_NAME], 1)]
```

Some of the brand names look like they are of the same brands - such as RED and RRD, which are both Red Rock Deli chips. Let's combine these together.

## Checking brands

```
transactionData[, .N, by = BRAND][order(-N)]
```

```
##      BRAND      N
##      <char> <int>
##  1:  Kettle 41288
##  2:   Smiths 27390
##  3: Pringles 25102
##  4:  Doritos 22041
##  5:    Thins 14075
##  6:     RRD 11894
##  7: Infuzions 11057
##  8:      WW 10320
##  9:     Cobs  9693
## 10: Tostitos  9471
## 11: Twisties  9454
## 12: Tyrrells  6442
## 13:   Grain  6272
## 14:  Natural  6050
## 15: Cheezels  4603
## 16:     CCs  4551
## 17:     Red  4427
## 18:  Dorito  3183
## 19:  Infzns  3144
## 20:   Smith  2963
## 21:  Cheetos  2927
## 22:   Snbts  1576
## 23:   Burger  1564
## 24: Woolworths 1516
## 25:   GrnWves 1468
## 26:  Sunbites 1432
## 27:     NCC  1419
## 28:   French 1418
##      BRAND      N
```

## Cleaning BRAND

```
transactionData[toupper(BRAND) == "RED", BRAND := "RRD"]
transactionData[BRAND == "Dorito", BRAND := "Doritos"]
transactionData[BRAND == "Smith", BRAND := "Smiths"]
```

```
transactionData[BRAND == "WW", BRAND := "Woolworths"]
transactionData[BRAND == "Grain", BRAND := "GrnWves"]
transactionData[BRAND == "Snbts", BRAND := "Sunbites"]
transactionData[BRAND == "Infzns", BRAND := "Infuzions"]
transactionData[BRAND == "NCC", BRAND := "Natural"]
```

Rechecking brands

```
transactionData[, .N, by = BRAND][order(-N)]
```

```
##      BRAND      N
##      <char> <int>
##  1:   Kettle 41288
##  2:   Smiths 30353
##  3:   Doritos 25224
##  4:  Pringles 25102
##  5:      RRD 16321
##  6: Infuzions 14201
##  7:     Thins 14075
##  8: Woolworths 11836
##  9:      Cobs  9693
## 10:  Tostitos  9471
## 11:  Twisties  9454
## 12:   GrnWves  7740
## 13:   Natural  7469
## 14:  Tyrrells  6442
## 15:  Cheezels  4603
## 16:      CCs  4551
## 17:  Sunbites  3008
## 18:   Cheetos  2927
## 19:   Burger  1564
## 20:   French  1418
##      BRAND      N
```

Finally, we have 20 different brands since 8 of our rows that had similar brands have been merged.

## Examining customer data

```
summary(customerData)
```

```
##  LYLTY_CARD_NBR      LIFESTAGE      PREMIUM_CUSTOMER
##  Min.   :   1000  Length:72637      Length:72637
##  1st Qu.: 66202   Class :character  Class :character
##  Median :134040   Mode  :character  Mode  :character
##  Mean   :136186
##  3rd Qu.:203375
##  Max.   :2373711
```

We can see that the loyalty card number is a numeric vector while lifestage and premium\_customer are character vectors. Next, I will join the transaction and customer data sets together.

## Merge data sets

```
data <- merge(transactionData, customerData, all.x = TRUE)
```

As the number of rows in data is the same as that of transactionData, we can be sure that no duplicates were created. This is because we created data by setting all.x = TRUE (in other words, a left join) which means take all the rows in transactionData and find rows with matching values in shared columns and then joining the details in these rows to the x or the first mentioned table.

## checking for nulls

```
colSums(is.na(data))
```

```
##  LYLTY_CARD_NBR      DATE      STORE_NBR      TXN_ID
##           0           0           0           0
##      PROD_NBR      PROD_NAME      PROD_QTY      TOT_SALES
##           0           0           0           0
##      PACK_SIZE      BRAND      LIFESTAGE PREMIUM_CUSTOMER
##           0           0           0           0
```

There are no nulls. So all our customers in transaction data has been accounted for in the customer dataset.

## Writing data as a csv

```
fwrite(data, paste0(filePath, "QVI_data.csv"))
```

## Data analysis on customer segment

### 1.Total Sales by LIFESTAGE and PREMIUM\_CUSTOMER

```
total_sales <- data %>%
  group_by(LIFESTAGE, PREMIUM_CUSTOMER) %>%
  summarise(SALSE = sum(TOT_SALES), .groups = "keep")
total_sales
```

```
## # A tibble: 21 x 3
## # Groups:   LIFESTAGE, PREMIUM_CUSTOMER [21]
##   LIFESTAGE      PREMIUM_CUSTOMER  SALSE
##   <chr>      <chr>      <dbl>
## 1 MIDGE SINGLES/COUPLES Budget      33346.
## 2 MIDGE SINGLES/COUPLES Mainstream  84734.
## 3 MIDGE SINGLES/COUPLES Premium      54444.
## 4 NEW FAMILIES      Budget      20607.
## 5 NEW FAMILIES      Mainstream  15980.
## 6 NEW FAMILIES      Premium      10761.
## 7 OLDER FAMILIES      Budget     156864.
## 8 OLDER FAMILIES      Mainstream   96414.
## 9 OLDER FAMILIES      Premium      75243.
## 10 OLDER SINGLES/COUPLES Budget     127834.
## # i 11 more rows
```

Plot for salse

```
ggplot(total_sales, aes(LIFESTAGE, SALSE, fill=PREMIUM_CUSTOMER)) +
  geom_bar(stat="identity", position=position_dodge()) +
  theme(axis.text.x = element_text(angle = 90)) +
  labs(title="Total Sales by Lifestage and Premium customer") +
  scale_fill_brewer(palette = "Set2")
```



Sales are coming mainly from Budget - older families, Mainstream - young singles/couples, and Mainstream - retirees

## 2.Number of customers by LIFESTAGE and PREMIUM\_CUSTOMER


```
customer <- data %>%
  group_by(LIFESTAGE, PREMIUM_CUSTOMER) %>%
  summarise(CUSTOMER = uniqueN(LYLT_CARD_NBR), .groups = "keep")
customer
```

```
## # A tibble: 21 x 3
## # Groups:   LIFESTAGE, PREMIUM_CUSTOMER [21]
##   LIFESTAGE          PREMIUM_CUSTOMER CUSTOMER
##   <chr>             <chr>             <int>
## 1 MIDAGE SINGLES/COUPLES Budget             1474
## 2 MIDAGE SINGLES/COUPLES Mainstream           3298
## 3 MIDAGE SINGLES/COUPLES Premium             2369
## 4 NEW FAMILIES          Budget             1087
## 5 NEW FAMILIES          Mainstream             830
```

```
## 6 NEW FAMILIES           Premium           575
## 7 OLDER FAMILIES         Budget           4611
## 8 OLDER FAMILIES         Mainstream        2788
## 9 OLDER FAMILIES         Premium          2231
## 10 OLDER SINGLES/COUPLES Budget          4849
## # i 11 more rows
```

Plot for number of customers

```
ggplot(customer, aes(LIFESTAGE, CUSTOMER, fill=PREMIUM_CUSTOMER)) +
  geom_bar(stat="identity", position=position_dodge()) +
  theme(axis.text.x = element_text(angle = 90)) +
  labs(title="Number of Customers by Lifestage and Premium customer") +
  scale_fill_brewer(palette = "Set2")
```



There are more Mainstream - young singles/couples and Mainstream - retirees who buy chips. This contributes to there being more sales to these customer segments but this is not a major driver for the Budget - Older families segment.

Higher sales may also be driven by more units of chips being bought per customer. Let's have a look at average number of units per customer by LIFESTAGE and PREMIUM\_CUSTOMER.

### 3.Average number of units per customer by LIFESTAGE and PREMIUM\_CUSTOMER

```
avg_unit <- data %>%
  group_by(LIFESTAGE, PREMIUM_CUSTOMER) %>%
```



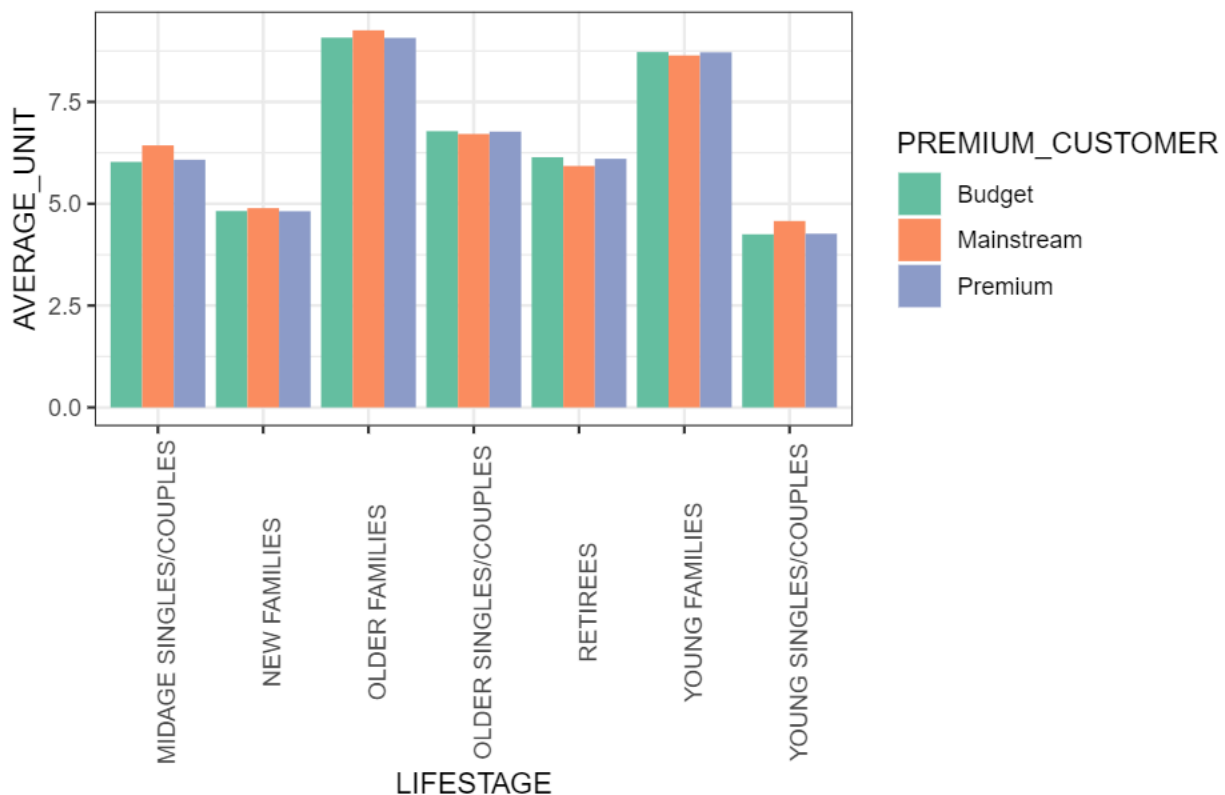
```
summarise(AVERAGE_UNIT = sum(PROD_QTY)/uniqueN(LYLTY_CARD_NBR), .groups = "keep")
avg_unit
```

```
## # A tibble: 21 x 3
## # Groups:   LIFESTAGE, PREMIUM_CUSTOMER [21]
##   LIFESTAGE          PREMIUM_CUSTOMER AVERAGE_UNIT
##   <chr>            <chr>            <dbl>
## 1 MIDAGE SINGLES/COUPLES Budget          6.03
## 2 MIDAGE SINGLES/COUPLES Mainstream        6.43
## 3 MIDAGE SINGLES/COUPLES Premium          6.08
## 4 NEW FAMILIES       Budget          4.82
## 5 NEW FAMILIES       Mainstream        4.89
## 6 NEW FAMILIES       Premium          4.82
## 7 OLDER FAMILIES     Budget          9.08
## 8 OLDER FAMILIES     Mainstream        9.26
## 9 OLDER FAMILIES     Premium          9.07
## 10 OLDER SINGLES/COUPLES Budget          6.78
## # i 11 more rows
```

Plot for average number of units per customer

```
ggplot(avg_unit, aes(LIFESTAGE, AVERAGE_UNIT, fill=PREMIUM_CUSTOMER)) +
  geom_bar(stat="identity", position=position_dodge()) +
  theme(axis.text.x = element_text(angle = 90)) +
  labs(title="Units per customer by Lifestage and Premium customer") +
  scale_fill_brewer(palette = "Set2")
```

Units per customer by Lifestage and Premium customer



Older families and young families in general buy more chips per customer.

Let's also investigate the average price per unit chips bought for each customer segment as this is also a driver of total sales.

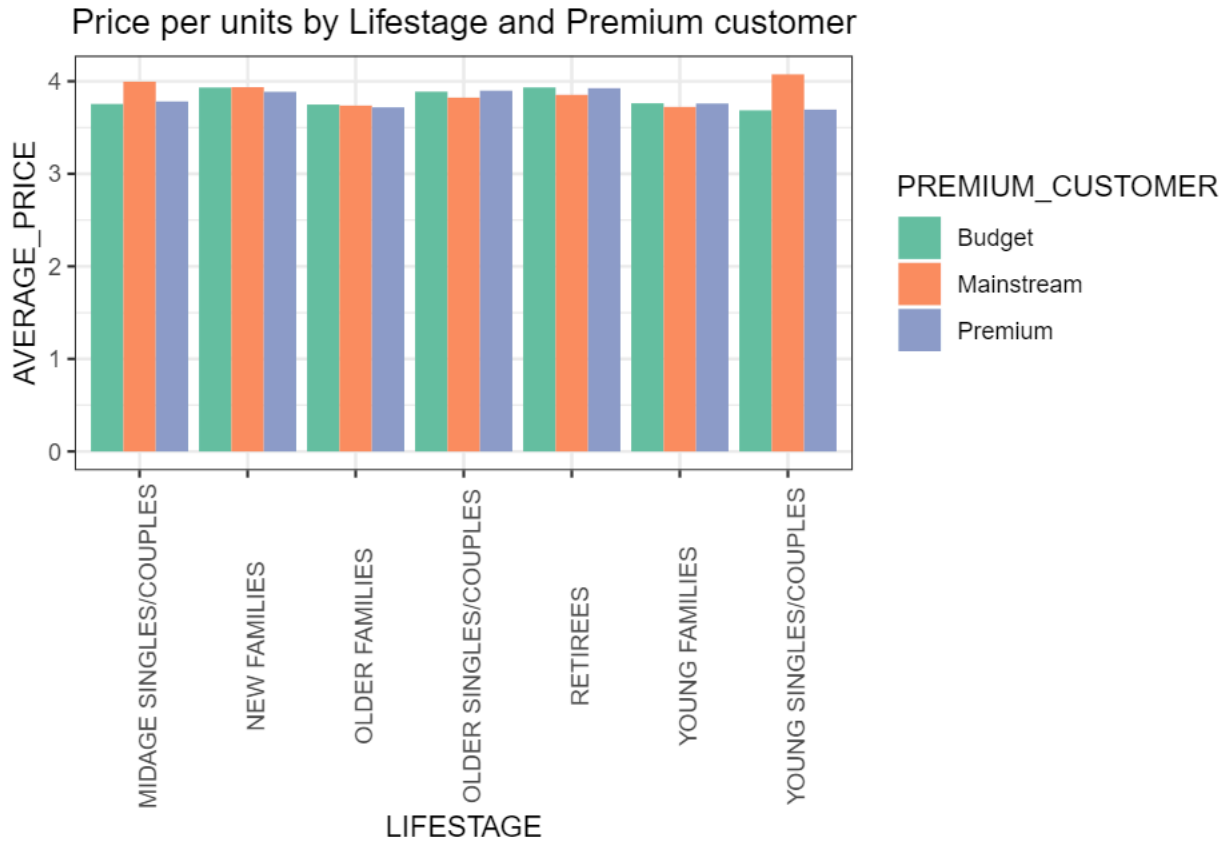
#### 4. Average price per unit by LIFESTAGE and PREMIUM\_CUSTOMER

```
avg_price <- data %>%
  group_by(LIFESTAGE, PREMIUM_CUSTOMER) %>%
  summarise(AVERAGE_PRICE = sum(TOT_SALES)/sum(PROD_QTY), .groups = "keep")
avg_price
```

```
## # A tibble: 21 x 3
## # Groups:   LIFESTAGE, PREMIUM_CUSTOMER [21]
##   LIFESTAGE          PREMIUM_CUSTOMER AVERAGE_PRICE
##   <chr>             <chr>             <dbl>
## 1 MIDAGE SINGLES/COUPLES Budget          3.75
## 2 MIDAGE SINGLES/COUPLES Mainstream        3.99
## 3 MIDAGE SINGLES/COUPLES Premium          3.78
## 4 NEW FAMILIES      Budget          3.93
## 5 NEW FAMILIES      Mainstream        3.94
## 6 NEW FAMILIES      Premium          3.89
## 7 OLDER FAMILIES    Budget          3.75
## 8 OLDER FAMILIES    Mainstream        3.74
## 9 OLDER FAMILIES    Premium          3.72
## 10 OLDER SINGLES/COUPLES Budget          3.89
## # i 11 more rows
```

Plot for average price per unit

```
ggplot(avg_price, aes(LIFESTAGE, AVERAGE_PRICE, fill=PREMIUM_CUSTOMER)) +
  geom_bar(stat="identity", position=position_dodge()) +
  theme(axis.text.x = element_text(angle = 90)) +
  labs(title="Price per units by Lifestage and Premium customer") +
  scale_fill_brewer(palette = "Set2")
```



Mainstream midage and young singles and couples are more willing to pay more per packet of chips compared to their budget and premium counterparts. This may be due to premium shoppers being more likely to buy healthy snacks and when they buy chips, this is mainly for entertainment purposes rather than their own consumption. This is also supported by there being fewer premium midage and young singles and couples buying chips compared to their mainstream counterparts.

## Conclusion

The data reveals valuable insights into customer segments, preferred brands, pack sizes, and spending trends.

- The purchasers of chips are coming mainly from Budget - older families, Mainstream - young singles/couples, and Mainstream - retirees shoppers. Among these groups, Mainstream - young singles/couples and retirees shoppers stand out for their high chip purchases due to their highest population.
- Kettle chips suggest an opportunity to capitalize on this by enhancing product visibility to attract more customers from this segment.
- The consistent preference for the 175-gram chip size followed by the 150-gram size across all customer segments indicates a strong market demand for these particular sizes.
- Sales peak just before Christmas, indicating a significant opportunity for increased revenue during this period. It's crucial to ensure sufficient stock levels to meet the heightened demand before Christmas, optimizing sales potential during this critical period.